Response to Justin's Questions

Which similarity metric seemed to work better? Why do you think that is?

Manhattan Distance barely worked. It would often find close to zero correct answers for a given file. Cosine distance and Euclidean distance had the same result and were both quite effective. There were no instances where at least a few matches could not be found. With a better vector model, the results could likely be even greater.

What input files did better than others?

Past tense (gram4) was far by the most effectively operating on. City and state (city-and-state) also yielded accurate solutions. Nationality adjective (gram6) also proved to be easier to analyze for the program. Familial Connections and superlatives tended to be harder for word_analogy to solve.

Did normalization help? In what cases?

The vectors were already normalized, therefore, normalization never made a difference in the output. There was no difference in content between output1x/eval.txt and output0x/eval.txt.